



Module std::io

1.0.0

[–][src]

[–] Traits, helpers, and type definitions for core I/O functionality.

The `std::io` module contains a number of common things you'll need when doing input and output. The most core part of this module is the `Read` and `Write` traits, which provide the most general interface for reading and writing input and output.

Read and Write

Because they are traits, `Read` and `Write` are implemented by a number of other types, and you can implement them for your types too. As such, you'll see a few different types of I/O throughout the documentation in this module: `Files`, `TcpStreams`, and sometimes even `Vec<T>`s. For example, `Read` adds a `read` method, which we can use on `Files`:

```
use std::io;
use std::io::prelude::*;
use std::fs::File;

fn main() -> io::Result<()> {
    let mut f = File::open("foo.txt")?;
    let mut buffer = [0; 10];

    // read up to 10 bytes
    let n = f.read(&mut buffer)?;

    println!("The bytes: {:?}", &buffer[..n]);
    Ok(())
}
```

Run

`Read` and `Write` are so important, implementors of the two traits have a nickname: readers and writers. So you'll sometimes see 'a reader' instead of 'a type that implements the `Read` trait'. Much easier!

Seek and BufRead

Beyond that, there are two important traits that are provided: `Seek` and `BufRead`. Both of these build on top of a reader to control how the reading happens. `Seek` lets you control where the next byte is coming from:

```
use std::io;
use std::io::prelude::*;
use std::io::SeekFrom;
use std::fs::File;

fn main() -> io::Result<()> {
    let mut f = File::open("foo.txt")?;
    let mut buffer = [0; 10];

    // skip to the last 10 bytes of the file
    f.seek(SeekFrom::End(-10))?;

    // read up to 10 bytes
    let n = f.read(&mut buffer)?;

    println!("The bytes: {:?}", &buffer[..n]);
    Ok(())
}
```

Run

`BufRead` uses an internal buffer to provide a number of other ways to read, but to show it off, we'll need to talk about buffers in general. Keep reading!

BufReader and BufWriter

Byte-based interfaces are unwieldy and can be inefficient, as we'd need to be making near-constant calls to the operating system. To help with this, `std::io` comes with two structs, `BufReader` and `BufWriter`, which wrap readers and writers. The wrapper uses a buffer, reducing the number of calls and providing nicer methods for accessing exactly what you want.

For example, `BufReader` works with the `BufRead` trait to add extra methods to any reader:

```
use std::io;
use std::io::prelude::*;
use std::io::BufReader;
use std::fs::File;

fn main() -> io::Result<()> {
    let f = File::open("foo.txt")?;
    let mut reader = BufReader::new(f);
    let mut buffer = String::new();

    // read a line into buffer
    reader.read_line(&mut buffer)?;

    println!("{}", buffer);
    Ok(())
}
```

Run

`BufWriter` doesn't add any new ways of writing; it just buffers every call to `write`:

```
use std::io;
use std::io::prelude::*;
use std::io::BufWriter;
use std::fs::File;

fn main() -> io::Result<()> {
    let f = File::create("foo.txt")?;
    {
        let mut writer = BufWriter::new(f);

        // write a byte to the buffer
        writer.write(&[42])?;

    } // the buffer is flushed once writer goes out of scope

    Ok(())
}
```

Run

Standard input and output

A very common source of input is standard input:

```
use std::io;

fn main() -> io::Result<()> {
    let mut input = String::new();

    io::stdin().read_line(&mut input)?;

    println!("You typed: {}", input.trim());
    Ok(())
}
```

Run

Note that you cannot use the `? operator` in functions that do not return a `Result<T, E>`. Instead, you can call `.unwrap()` or `match` on the return value to catch any possible errors:

```
use std::io;

let mut input = String::new();

io::stdin().read_line(&mut input).unwrap();
```

Run

And a very common source of output is standard output:

```
use std::io;
use std::io::prelude::*;

fn main() -> io::Result<()> {
    io::stdout().write(&[42])?;
    Ok(())
}
```

Run

Of course, using `io::stdout` directly is less common than something like `println!`.

Iterator types

A large number of the structures provided by `std::io` are for various ways of iterating over I/O. For example, `Lines` is used to split over lines:

```
use std::io;
use std::io::prelude::*;
use std::io::BufReader;
use std::fs::File;

fn main() -> io::Result<()> {
    let f = File::open("foo.txt")?;
    let reader = BufReader::new(f);

    for line in reader.lines() {
        println!("{}", line?);
    }
    Ok(())
}
```

Run

Functions

There are a number of `functions` that offer access to various features. For example, we can use three of these functions to copy everything from standard input to standard output:

```
use std::io;

fn main() -> io::Result<()> {
    io::copy(&mut io::stdin(), &mut io::stdout())?;
    Ok(())
}
```

Run

io::Result

Last, but certainly not least, is `io::Result`. This type is used as the return type of many `std::io` functions that can cause an error, and can be returned from your own functions as well. Many of the examples in this module use the `? operator`:

```
use std::io;

fn read_input() -> io::Result<()> {
    let mut input = String::new();

    io::stdin().read_line(&mut input)?;

    println!("You typed: {}", input.trim());

    Ok(())
}
```

Run

The return type of `read_input()`, `io::Result<()>`, is a very common type for functions which don't have a 'real' return value, but do want to return errors if they happen. In this case, the only purpose of this function is to read the line and print it, so we use `()`.

Platform-specific behavior

Many I/O functions throughout the standard library are documented to indicate what various library or syscalls they are delegated to. This is done to help applications both understand what's happening under the hood as well as investigate any possibly unclear semantics. Note, however, that this is informative, not a binding contract. The implementation of many of these functions are subject to change over time and may call fewer or more syscalls/library functions.

Modules

`prelude` The I/O Prelude

Structs

<code>Initializer</code>	<div><div>Experimental</div>A type used to conditionally initialize buffers passed to <code>Read</code> methods.</div>
<code>BufReader</code>	The <code>BufReader<R></code> struct adds buffering to any reader.
<code>BufWriter</code>	Wraps a writer and buffers its output.
<code>Bytes</code>	An iterator over <code>u8</code> values of a reader.
<code>Chain</code>	Adaptor to chain together two readers.
<code>Cursor</code>	A <code>Cursor</code> wraps an in-memory buffer and provides it with a <code>Seek</code> implementation.
<code>Empty</code>	A reader which is always at EOF.
<code>Error</code>	The error type for I/O operations of the <code>Read</code> , <code>Write</code> , <code>Seek</code> , and associated traits.
<code>IntoInnnerError</code>	An error returned by <code>BufWriter::into_inner</code> which combines an error that happened while writing out the buffer, and the buffered writer object which may be used to recover from the condition.
<code>IoSlice</code>	A buffer type used with <code>Write::write_vectored</code> .
<code>IoSliceMut</code>	A buffer type used with <code>Read::read_vectored</code> .
<code>LineWriter</code>	Wraps a writer and buffers output to it, flushing whenever a newline (<code>0x0a</code> , <code>'\n'</code>) is detected.
<code>Lines</code>	An iterator over the lines of an instance of <code>BufRead</code> .
<code>Repeat</code>	A reader which yields one byte over and over and over and over and...
<code>Sink</code>	A writer which will move data into the void.
<code>Split</code>	An iterator over the contents of an instance of <code>BufRead</code> split on a particular byte.
<code>Stderr</code>	A handle to the standard error stream of a process.
<code>StderrLock</code>	A locked reference to the <code>Stderr</code> handle.
<code>Stdin</code>	A handle to the standard input stream of a process.
<code>StdinLock</code>	A locked reference to the <code>Stdin</code> handle.
<code>Stdout</code>	A handle to the global standard output stream of the current process.
<code>StdoutLock</code>	A locked reference to the <code>Stdout</code> handle.
<code>Take</code>	Reader adaptor which limits the bytes read from an underlying reader.

Enums

<code>ErrorKind</code>	A list specifying general categories of I/O error.
<code>SeekFrom</code>	Enumeration of possible methods to seek within an I/O object.

Traits

<code>BufRead</code>	A <code>BufRead</code> is a type of <code>Reader</code> which has an internal buffer, allowing it to perform extra ways of reading.
<code>Read</code>	The <code>Read</code> trait allows for reading bytes from a source.
<code>Seek</code>	The <code>Seek</code> trait provides a cursor which can be moved within a stream of bytes.
<code>Write</code>	A trait for objects which are byte-oriented sinks.

Functions

<code>copy</code>	Copies the entire contents of a reader into a writer.
<code>empty</code>	Constructs a new handle to an empty reader.
<code>repeat</code>	Creates an instance of a reader that infinitely repeats one byte.
<code>sink</code>	Creates an instance of a writer which will successfully consume all data.
<code>stderr</code>	Constructs a new handle to the standard error of the current process.
<code>stdin</code>	Constructs a new handle to the standard input of the current process.
<code>stdout</code>	Constructs a new handle to the standard output of the current process.

Type Definitions

`Result` A specialized `Result` type for I/O operations.

Module io

Modules

Structs

Enums

Traits

Functions

Type Definitions

std

Primitive Types

array

bool

char

f32

f64

fn

i128

i16

i32

i64

i8

isize

never

pointer

reference

slice

str

tuple

u128

u16

u32

u64

u8

unit

usize

Modules

alloc

any

array

ascii

backtrace

borrow

boxed

cell

char

clone

cmp

collections

convert

default

env

error

f32

f64

ffi

fmt

fs

future

hash

hint

i128

i16

i32

i64

i8

intrinsics

io

isize

iter

lazy

marker

mem

net

num

ops

option

os

panic

path

pin

prelude

primitive

process

ptr

raw

rc

result

slice

str

string

sync

task

thread

time

u128

u16

u32

u64

u8

usize

vec

Macros

asm

assert

assert_eq

assert_ne

cfg

column

compile_error

concat

concat_idents

dbg

debug_assert

debug_assert_eq

debug_assert_ne

env

eprint

eprintln

file

format

format_args

format_args_nl

global_asm

include

include_bytes

include_str

is_aarch64_feature_detected

is_arm_feature_detected

is_mips64_feature_detected

is_mips_feature_detected

is_powerpc64_feature_detected

is_powerpc_feature_detected

is_x86_feature_detected

line

llvm_asm

log_syntax

matches

module_path

option_env

panic

print

println

stringify

thread_local

todo

trace_macros

try

unimplemented

unreachable

vec

write

writeln

Keywords

Self

as

async

await

break

const

continue

crate

dyn

else

enum

extern

false

fn

for

if

impl

in

let

loop

match

mod

move

mut

pub

ref

return

self

static

struct

super

trait

true

type

union

unsafe

use

where

while